Replicating xRFM: Accurate, Scalable, and Interpretable Feature Learning Models for Tabular Data +bonus

COMP9417 Group Project

Xueyao Jiang z5581908 Chang Yang z5569348 Binyue Peng z5705242
Jiayu Zhang z5649264 Zirun Zhang z5617187

April 28, 2026

Contents

1 Introduction **1**

2 Methodology **1**
2.1 xRFM: Algorithm and Theoretical Background . . . 1
2.2 Experimental Design . . . 2
2.2.1 Datasets . . . 2
2.2.2 Preprocessing and Splitting . . . 2
2.2.3 Evaluation Metrics . . . 2
2.2.4 xRFM Configuration . . . 2
2.2.5 XGBoost and LightGBM Configuration . . . 3
2.2.6 MLP Configuration . . . 3

3 Results **3**
3.1 Predictive Performance . . . 3
3.2 Scalability Analysis . . . 4
3.3 Interpretability Comparison . . . 5

4 Discussion **7**
4.1 xRFM: When Feature Learning Helps and When It Does Not . . . 7
4.2 XGBoost and LightGBM: Performance and Efficiency . . . 7
4.3 MLP: Performance and Scalability Limitations . . . 7

5 Conclusion **7**

Appendix **9**

A Bonus: Residual-Weighted AGOP **9**

B Supplementary Figures **10**

References **11**

1 Introduction

Tabular data remains one of the most common data modalities in applied machine learning, appearing in domains such as materials science, healthcare, e-commerce, and human-resource analytics. Despite the rapid progress of deep learning on images and language, GBDTs such as XGBoost and LightGBM remain difficult to outperform on many tabular tasks [Grinsztajn et al., 2022, Holzmüller et al., 2024].

This project replicates and evaluates xRFM, a recent feature learning approach for tabular data based on the average gradient outer product (AGOP) [Beaglehole et al., 2025, Radhakrishnan et al., 2024]. Instead of relying only on fixed input features, xRFM uses gradient-based information from a learned prediction function to estimate directions in the input space that are most relevant to the target. This produces both a predictive model and an interpretable feature-relevance signal. The central question of our replication is therefore whether AGOP-driven feature learning can offer a competitive alternative to strong tabular baselines, especially when compared with XGBoost, LightGBM, and a multilayer perceptron (MLP).

We evaluate all models under a unified pipeline across five public tabular datasets covering regression and binary classification, large and small sample sizes, and mixed feature types (Diamonds, Superconductivity, Online Shoppers, Stroke, HR Attrition). Our main finding is that tree-based GBDTs lead overall, while xRFM shows selective strength on small-sample and fully numerical tasks. Section 2 describes methodology, Section 3 reports results, and Sections 4–5 discuss findings and conclusions.

2 Methodology

2.1 xRFM: Algorithm and Theoretical Background

xRFM is a tabular feature-learning model based on the Average Gradient Outer Product (AGOP) [Beaglehole et al., 2025, Radhakrishnan et al., 2024]. For a trained prediction function $f : \mathbb{R}^d \to \mathbb{R}$, the AGOP is defined as

$$G_f = \mathbb{E}_x\left[\nabla f(x) \nabla f(x)^\top\right] \in \mathbb{R}^{d \times d}. \tag{1}$$

The diagonal entry $[G_f]_{ii} = \mathbb{E}_x[(\partial f/\partial x_i)^2]$ is the expected squared partial derivative with respect to feature i: a large value indicates rapid change in the model output along coordinate i (important feature); a small value indicates near-insensitivity. The full G_f supports two interpretability strategies: its *diagonal* gives coordinate-wise importance scores for each feature, while its *top eigenvectors* identify multivariate directions in input space along which the model varies most, potentially capturing signal invisible to any single coordinate. We use diagonal AGOP in our experiments as it is cheaper to store and sufficient for feature-level rankings.

Compared with alternative methods, AGOP is supervised and model-specific: it reflects the gradient geometry of the fitted function, unlike PCA which only uses input covariance [Pearson, 1901]. Mutual information measures marginal feature–target dependence but not the model's multivariate sensitivity [Kraskov et al., 2004]. Permutation importance is model-agnostic but measures post-hoc performance degradation rather than learned geometry [Fisher et al., 2019]. AGOP is extracted as a by-product of training and directly summarises which directions the learned model relies on.

Training pipeline. xRFM builds a binary tree over the training data. At each node, a base model f is fitted to the local data partition and its AGOP G_f is computed. The leading eigenvector v_1 of G_f defines the split direction; the threshold t^* minimises weighted target

variance along the projection $v_1^\top x$. Each leaf fits a local random Fourier feature kernel machine with the feature metric reweighted by G_f. For inference on a new point x^*, the tree is traversed via $v_k^\top x^* \leq t_k$ at each internal node, and the leaf model produces the final prediction. The dominant training cost is AGOP computation: $O(nd)$ per node for a linear approximation, giving total complexity $O(nd \cdot 2^D)$ for tree depth D [Beaglehole et al., 2025]. This addresses the fundamental limitation of standard kernel machines on tabular data — that a global Euclidean metric treats all features equally — by adaptively reweighting feature space at each tree node.

2.2 Experimental Design

2.2.1 Datasets

Five public tabular datasets (Table 1) are used, spanning regression and binary classification, small and large sample sizes, and mixed feature types [Hamidieh, 2018, Sakar et al., 2019, fedesoriano, 2021].

Table 1: Dataset statistics. n_{train}, n_{val}, and n_{test} denote the sizes of the training, validation, and test splits respectively; d is the number of features after preprocessing.

Dataset	Task	Target	n_{train}	n_{val}	n_{test}	d
Diamonds	Regression	price	32,365	10,789	10,789	26
Superconductivity	Regression	critical_temp	12,757	4,253	4,253	81
Online Shoppers	Classification	Revenue	7,398	2,466	2,466	27
Stroke	Classification	stroke	3,066	1,022	1,022	21
HR Attrition	Classification	Attrition	882	294	294	51

2.2.2 Preprocessing and Splitting

All experiments use seed 42 and a 60/20/20 split. Numerical features are z-score normalised and categorical features are one-hot encoded; missing values are median/constant imputed. The same encoded matrix is used across all four models.

2.2.3 Evaluation Metrics

Regression tasks are evaluated with RMSE; binary classification tasks with accuracy and AUC-ROC. Training and inference time are also recorded.

2.2.4 xRFM Configuration

The xRFM experiments use the Python `xrfm` package through the shared project training pipeline. The model is initialised as `xRFM(use_diag=True, device=DEVICE)`, where `DEVICE` is CUDA if available and CPU otherwise. On GPU, the training script sets the number of fit iterations to 3 and `M_batch_size=128` to control memory cost. Feature importance is extracted by reading the learned AGOP matrix or diagonal values from the trained xRFM model and ranking features by their AGOP-derived scores.

Table 2: xRFM configuration used in the experiments.

Component	Setting
Model	`xRFM`
AGOP mode	Diagonal AGOP, `use_diag=True`
Device	CUDA if available; otherwise CPU
Random seed	42
GPU fit iterations	3
GPU AGOP batch size	128
Data split	60/20/20 train/validation/test
Regression metric	RMSE
Classification metrics	Accuracy and AUC-ROC
Interpretability output	Top features ranked by AGOP diagonal score

2.2.5 XGBoost and LightGBM Configuration

XGBoost and LightGBM are trained through the scikit-learn API using strong default settings from recent tabular benchmarking work [Chen and Guestrin, 2016, Ke et al., 2017, Holzmüller et al., 2024]. Both use 1,000 maximum boosting rounds, learning rate $\eta = 0.05$, row and column subsampling of 0.8, and early stopping with patience 20 on the validation set. XGBoost uses maximum tree depth 6, while LightGBM uses `num_leaves=31`.

2.2.6 MLP Configuration

The MLP baseline uses a two-hidden-layer network with hidden sizes $(128, 64)$ and ReLU activation. It is trained with Adam, learning rate 0.001, batch size 256, maximum 500 iterations, and early stopping with patience 20. Input features are standardised before training. For classification tasks, predicted probabilities are used for AUC-ROC calculation.

3 Results

This section compares xRFM with XGBoost, LightGBM, and MLP in terms of predictive performance, computational efficiency, scalability, and interpretability. All models are evaluated using the same preprocessing pipeline, train-validation-test split, and evaluation metrics described in Section 2.

3.1 Predictive Performance

Table 3 reports the final held-out test performance across all five datasets. RMSE is used for regression tasks, while accuracy and AUC-ROC are used for binary classification tasks. Bold values indicate the best result for each metric.

Regression tasks. LightGBM achieves the lowest RMSE on both regression datasets. xRFM is competitive on Superconductivity, where its RMSE of 9.454 is close to LightGBM's 9.209 and XGBoost's 9.302. However, xRFM performs much worse on Diamonds, with RMSE 1496.72 compared with 535.99 for LightGBM and 544.68 for XGBoost. This suggests that AGOP-based feature learning does not automatically outperform tree ensembles on large mixed-type tabular regression.

Classification tasks. For Online Shoppers and Stroke, XGBoost and LightGBM obtain the strongest AUC-ROC scores. xRFM remains reasonably competitive on Online Shoppers, but

Table 3: Test-set performance of all models. **Bold** indicates the best result per dataset and metric. RMSE is lower-better; Accuracy and AUC are higher-better.

Dataset	Metric	Model			
		xRFM	XGBoost	LightGBM	MLP
Diamonds	RMSE ↓	1496.72	544.68	**535.99**	734.69
Superconductivity	RMSE ↓	9.454	9.302	**9.209**	11.066
Online Shoppers	Accuracy ↑	0.8897	0.9031	**0.9039**	0.8946
	AUC ↑	0.9018	**0.9306**	0.9306	0.9040
Stroke	Accuracy ↑	0.9413	0.9491	0.9501	**0.9511**
	AUC ↑	0.7605	**0.8388**	0.8231	0.2690
HR Attrition	Accuracy ↑	**0.8673**	0.8571	0.8469	0.8639
	AUC ↑	0.7631	0.7552	**0.7886**	0.7512

its AUC is below the boosted-tree models. xRFM's strongest result is on HR Attrition, where it achieves the highest accuracy of 0.8673, suggesting a modest advantage in small-sample classification. MLP obtains high accuracy on Stroke but very low AUC, showing that accuracy alone can be misleading under severe class imbalance.

Table 4 reports training time and per-sample inference time for all models. The timing results show that xRFM can be efficient on smaller datasets, but becomes very expensive on Diamonds. In contrast, LightGBM and XGBoost remain consistently fast.

Table 4: Training time (s) and inference time per sample (μs) for all models. Each cell shows Train / Infer.

Dataset	xRFM	XGBoost	LightGBM	MLP
Diamonds	246.5 s / 69.8 μs	1.4 s / 1.9 μs	0.8 s / 1.9 μs	26.8 s / 0.9 μs
Superconductivity	1.6 s / 1.0 μs	3.4 s / 2.1 μs	1.9 s / 7.6 μs	34.3 s / 0.8 μs
Online Shoppers	0.7 s / 0.8 μs	0.2 s / 1.4 μs	0.3 s / 1.6 μs	2.1 s / 0.8 μs
Stroke	0.3 s / 1.1 μs	0.1 s / 1.1 μs	0.1 s / 2.0 μs	0.4 s / 1.0 μs
HR Attrition	0.1 s / 3.2 μs	0.1 s / 3.0 μs	0.1 s / 7.5 μs	0.4 s / 2.2 μs

3.2 Scalability Analysis

To evaluate scalability, we run subsampling experiments on Diamonds and Superconductivity. Table 5 reports RMSE and training time at different training sizes. Figure 1 visualises the Diamonds scalability curves.

Table 5: Scalability experiment: test RMSE and training time (seconds) as a function of training set size. Dash (–) indicates configurations not run.

Dataset	n_{train}	xRFM		XGBoost		LightGBM		MLP	
		RMSE	Time (s)	RMSE	Time (s)	RMSE	Time (s)	RMSE	Time (s)
Diamonds	1,000	1657.0	0.77	846.7	0.80	1003.4	0.13	1081.0	2.27
	5,000	1655.9	11.03	628.9	0.39	622.7	0.11	901.1	10.12
	10,000	1577.6	39.31	587.8	0.49	592.3	0.15	888.9	10.33
	30,000	1487.5	226.44	542.7	0.59	535.3	0.29	743.5	22.23
Supercond.	1,000	14.07	0.08	13.76	0.39	13.91	0.15	15.39	2.52
	5,000	11.30	0.47	10.81	1.00	10.71	0.46	13.05	8.58
	10,000	9.81	1.07	9.51	1.73	9.50	1.10	11.22	28.35
	12,757	9.46	1.56	9.28	1.64	9.21	1.04	–	–

On Diamonds, xRFM shows a clear scalability limitation. Its training time increases from 0.77s to 226.44s as n_{train} grows from 1,000 to 30,000, while RMSE only improves from 1657.0 to 1487.5. In contrast, LightGBM reaches 535.3 RMSE at 30,000 training samples while training in less than one second. On Superconductivity, xRFM scales more favourably: its RMSE improves from 14.07 to 9.46 and remains close to XGBoost and LightGBM at full size. This contrast suggests that xRFM is better suited to fully numerical, high-dimensional regression than to large mixed-type datasets with strong threshold and categorical effects.

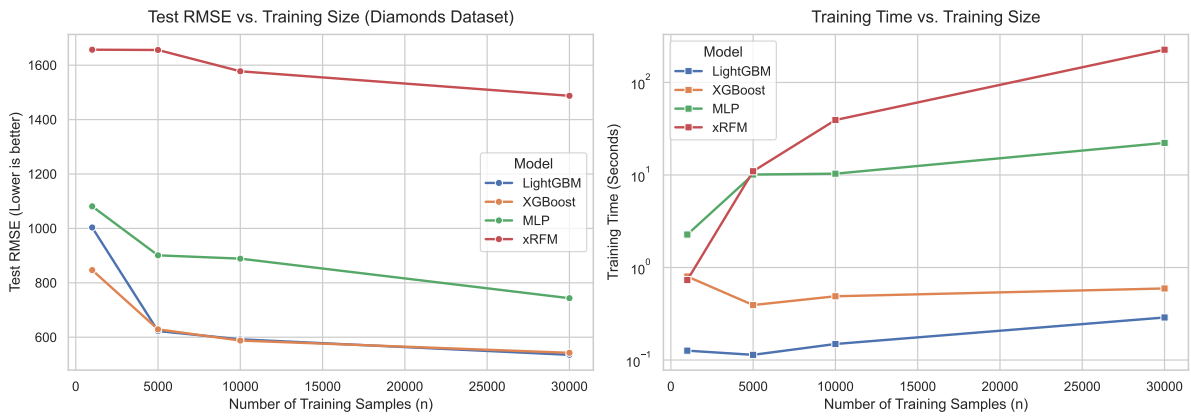

Figure 1: Scalability analysis on the Diamonds dataset, comparing test RMSE (left) and training time (right) across varying training set sizes.

3.3 Interpretability Comparison

Table 6 lists the top-five AGOP-ranked features extracted from each trained xRFM model. These results show that the AGOP diagonal often highlights meaningful variables: physical size for Diamonds, behavioural value for Online Shoppers, health risk factors for Stroke, and work-intensity indicators for HR Attrition.

Table 6: xRFM AGOP-diagonal feature importance extracted from the trained models.

Dataset	Top AGOP-ranked features, from 1st to 5th
Diamonds	y, x, carat, z, depth
Superconductivity	number_of_elements, mean_atomic_mass, wtd_mean_atomic_mass, gmean_atomic_mass, wtd_gmean_atomic_mass
Online Shoppers	PageValues, ExitRates, ProductRelated, ProductRelated_Duration, BounceRates
Stroke	avg_glucose_level, age, bmi, work_type_Self-employed, smoking_status_smokes
HR Attrition	OverTime_No, OverTime_Yes, BusinessTravel_Travel_Frequently, TotalWorkingYears, BusinessTravel_Non-Travel

To answer whether AGOP agrees with other feature-importance methods, we perform a detailed comparison on the Diamonds dataset. Table 7 compares AGOP diagonal importance against PCA loadings, mutual information, and permutation importance on the same data.

Table 7: Feature-importance comparison on the Diamonds dataset.

Method	Top ranked features
AGOP diagonal	y, x, carat, z, depth
PCA loadings	x, y, z, carat, table
Mutual information	carat, x, y, z, clarity_SI1
Permutation importance	carat, y, x, z, color_I

Figure 2 visualises the same comparison. It shows that AGOP, mutual information, and permutation importance agree on the importance of physical size-related variables, while permutation importance gives relatively more weight to some categorical quality variables.

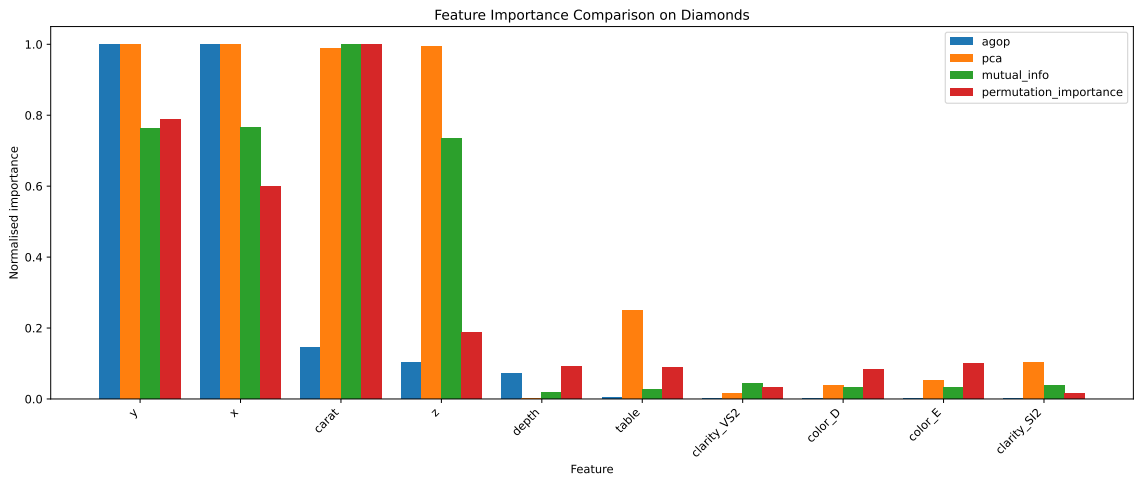

Figure 2: Feature-importance comparison on the Diamonds dataset using AGOP diagonal scores, PCA loadings, mutual information, and permutation importance.

The empirical comparison partly supports the theoretical argument made in Section 2.1. AGOP is supervised and model-specific, so it should identify features that the learned xRFM function is sensitive to. On Diamonds, AGOP ranks y, x, carat, and z among the most important features,

which agrees with mutual information and permutation importance that physical size and carat weight are central predictors of diamond price. PCA also ranks `x`, `y`, `z`, and `carat` highly, but this is expected because these variables dominate input variance rather than necessarily target relevance.

The main divergence is that permutation importance assigns additional importance to categorical quality features such as `color_I`, while AGOP focuses more heavily on continuous size-related coordinates. This difference is informative: AGOP reflects the sensitivity structure of the fitted xRFM model, whereas permutation importance directly measures test-performance degradation after feature corruption. A surprising finding is that AGOP produces intuitively meaningful feature rankings even though xRFM performs poorly on Diamonds compared with XGBoost and LightGBM. This suggests that interpretability and predictive accuracy are not equivalent. AGOP can identify sensible prediction-relevant directions, but the xRFM prediction pipeline may still be less effective than boosted trees at modelling sharp axis-aligned thresholds and categorical interactions in this dataset.

4 Discussion

4.1 xRFM: When Feature Learning Helps and When It Does Not

xRFM's strongest result is HR Attrition ($n = 882$), where its structured feature-learning prior may help in low-data settings. However, on Diamonds it obtains RMSE 1496.72 versus 535.99 for LightGBM — nearly three times worse. Diamonds contains strong axis-aligned categorical effects (cut, colour, clarity) that boosted trees capture efficiently through split-based interactions [Chen and Guestrin, 2016, Ke et al., 2017], whereas AGOP gradient reweighting is less effective here. The scalability experiment confirms that adding more Diamonds data does not close this gap. In contrast, on the fully numerical Superconductivity dataset xRFM approaches the tree baselines, suggesting it is better suited to continuous high-dimensional inputs. Future work should target categorical handling, more efficient AGOP approximation, and dataset-specific tuning.

4.2 XGBoost and LightGBM: Performance and Efficiency

XGBoost and LightGBM deliver competitive or superior performance across all five datasets while remaining the fastest models to train. On Diamonds, both GBDTs outperform xRFM by a large margin in RMSE (544.68 and 535.99 versus 1496.72), consistent with the strong axis-aligned inductive bias of tree-based splits [Grinsztajn et al., 2022]. On HR Attrition ($n = 882$, $d = 51$), xRFM marginally surpasses XGBoost in accuracy (0.8673 versus 0.8571), suggesting that kernel-style feature learning offers a modest advantage only in a restricted small-data regime. LightGBM is the fastest model at every large Diamonds scale, with near-linear growth in training time (Table 5).

4.3 MLP: Performance and Scalability Limitations

MLP underperforms both GBDTs on regression (RMSE 743.5 vs. 535.3 on Diamonds at $n =$ 30,000) and shows unstable classification results, most notably AUC 0.269 on Stroke. This reflects sensitivity to class imbalance and the limited alignment between global gradient-based optimisation and the piecewise structure of tabular features [Holzmüller et al., 2024].

5 Conclusion

We replicated xRFM and compared it with XGBoost, LightGBM, and MLP across five tabular datasets under a unified experimental pipeline. Tree-based GBDTs remain the strongest general-purpose baselines: LightGBM and XGBoost lead on accuracy and training speed across nearly

all settings. xRFM shows selective promise — competitive on Superconductivity and best accuracy on HR Attrition — but underperforms substantially on Diamonds and scales poorly there, indicating that AGOP feature learning is not universally superior to tree splitting on large mixed-type data. AGOP-diagonal rankings are broadly interpretable and align with domain knowledge, though meaningful rankings do not guarantee predictive strength. Future work should address categorical variable handling, AGOP approximation efficiency, and hyperparameter tuning for diverse tabular settings.

A Bonus: Residual-Weighted AGOP

(i) Justification. Standard AGOP weights all samples equally ($G_f = \frac{1}{n}\sum_i \nabla f(x_i)\nabla f(x_i)^\top$). Points with large residuals $r_i = y_i - f(x_i)$ indicate regions where the model fails. Upweighting by $\varphi(r_i) = r_i^2$ shifts the AGOP toward directions of high model error, providing a more focused signal for guided tree splits.

(ii) Implementation. The residual-weighted AGOP is defined as

$$\mathrm{AGOP}_{\mathrm{res}}(f) = \frac{\displaystyle\sum_{i=1}^{n} r_i^2\, \nabla f(x_i)\nabla f(x_i)^\top}{\displaystyle\sum_{i=1}^{n} r_i^2}. \tag{2}$$

Gradients are estimated by forward finite differences ($\varepsilon = 10^{-3}$) on a degree-2 Ridge regression model. The core implementation is two lines:

```
def agop(G, residuals=None):
    w = np.ones(len(G))/len(G) if residuals is None else residuals**2
    w /= w.sum() + 1e-12          # normalise; 1e-12 guards div-by-zero
    return np.einsum('i,ij,ik->jk', w, G, G)
```

On a single-regime dataset (n=200, d=6, $y = 4x_0 + 3\sin(2x_1) + 0.5\varepsilon$), both AGOP variants recover $\{x_0, x_1\}$ as the top features with cosine similarity 0.9994, confirming correct implementation.

(iii) Disagreement. We construct a two-regime dataset (n=600, d=6) designed to expose the difference between the two AGOP variants:

- Regime A ($x_5 < 0$): $y = 5x_0 + 0.05\varepsilon$ — well-fitted, residual RMSE 2.20.
- Regime B ($x_5 \geq 0$): $y = 5x_2 + 5.0\varepsilon$ — poorly fitted, residual RMSE 5.10.

A global Poly-2 Ridge model (α=0.5) is fitted to the combined data. Standard AGOP identifies x_0 (loading 0.863) as the leading feature: the well-fitted regime A dominates because its gradients are consistent. Residual-weighted AGOP identifies x_2 (loading 0.682): the large residuals in regime B inflate the weights, shifting the leading direction toward the hard-to-fit signal. The cosine similarity between the two leading eigenvectors is 0.676, confirming substantial directional disagreement ($p = 1$).

(iv) Performance. Using the same two-regime DGP (n_{train}=500, n_{test}=200), we compare RMSE of a one-level tree split using each AGOP direction:

Method	Test RMSE
Global model (no split)	4.034
Standard AGOP split (along x_0)	4.151
Res-weighted AGOP split (along x_2)	4.173

Both single-level splits produce slightly higher RMSE than the global model, and neither variant achieves a clear advantage over the other. This suggests that in this controlled setting, residual weighting's main value is *interpretive*—correctly identifying where the model fails—rather than automatically improving a single-split predictor. Richer tree structures or iterative reweighting may be needed to realise a performance benefit.

B Supplementary Figures

AGOP Feature Importance

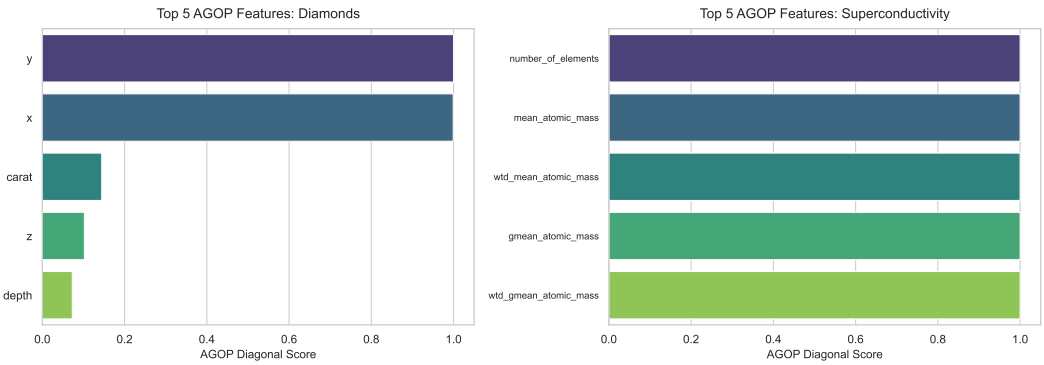

Figure 3: Top-five AGOP diagonal feature importance scores for Diamonds and Superconductivity.

Superconductivity Scalability

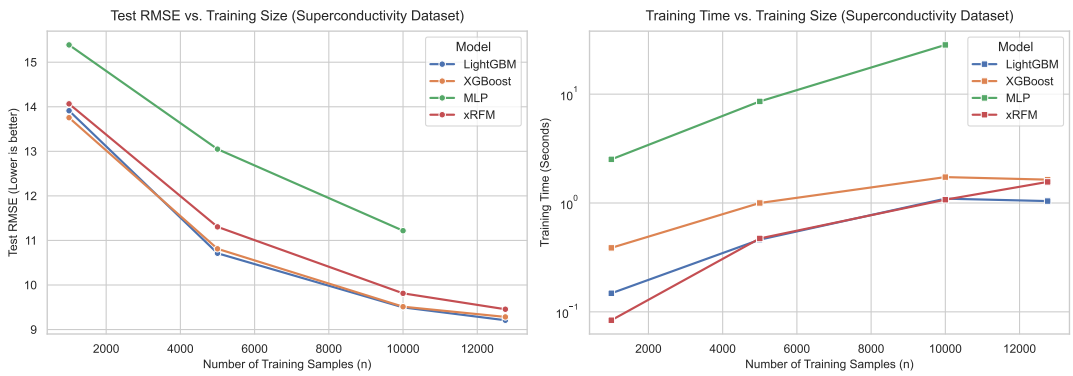

Figure 4: Scalability on Superconductivity: test RMSE and training time vs. training-set size.

Training-Time and Accuracy Trade-off

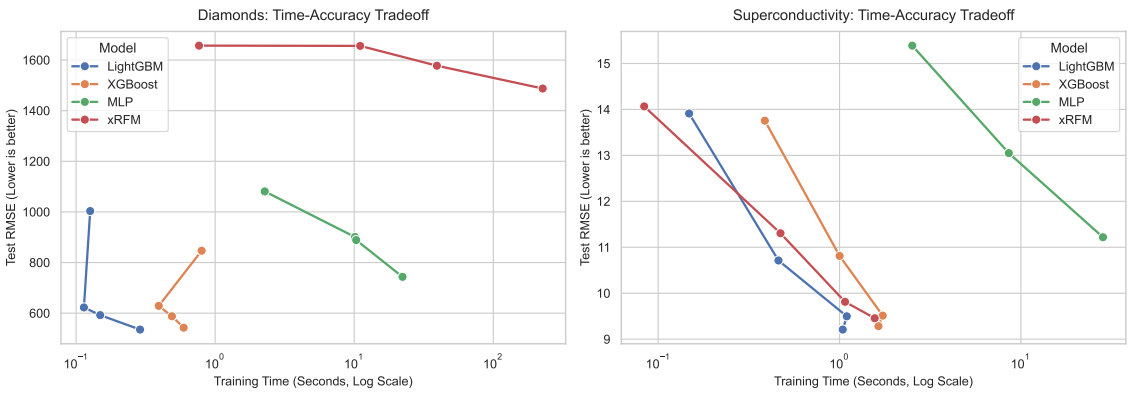

Figure 5: Training-time vs. predictive-error trade-off on Diamonds and Superconductivity.

References

Daniel Beaglehole, David Holzmüller, Adityanarayanan Radhakrishnan, and Mikhail Belkin. xrfm: Accurate, scalable, and interpretable feature learning models for tabular data. *arXiv preprint arXiv:2508.10053*, 2025.

Tianqi Chen and Carlos Guestrin. Xgboost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016.

fedesoriano. Stroke prediction dataset, 2021. URL https://www.kaggle.com/fedesoriano/stroke-prediction-dataset.

Aaron Fisher, Cynthia Rudin, and Francesca Dominici. All models are wrong, but many are useful: Learning a variable's importance by studying an entire class of prediction models simultaneously. *Journal of Machine Learning Research*, 20(177):1–81, 2019.

Léo Grinsztajn, Edouard Oyallon, and Gael Varoquaux. Why do tree-based models still outperform deep learning on typical tabular data? In *Advances in Neural Information Processing Systems*, volume 35, pages 507–520, 2022.

Kam Hamidieh. A data-driven statistical model for predicting the critical temperature of a superconductor. *Computational Materials Science*, 154:346–354, 2018.

David Holzmüller, Léo Grinsztajn, and Ingo Steinwart. Better by default: Strong pre-tuned mlps and boosted trees on tabular data. In *Advances in Neural Information Processing Systems*, volume 37, pages 26577–26658, 2024.

Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. Lightgbm: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

Alexander Kraskov, Harald Stögbauer, and Peter Grassberger. Estimating mutual information. *Physical Review E*, 69(6):066138, 2004.

Karl Pearson. On lines and planes of closest fit to systems of points in space. *The London, Edinburgh, and Dublin Philosophical Magazine and Journal of Science*, 2(11):559–572, 1901.

Adityanarayanan Radhakrishnan, Daniel Beaglehole, Parthe Pandit, and Mikhail Belkin. Mechanism for feature learning in neural networks and backpropagation-free machine learning models. *Science*, 383(6690):1461–1467, 2024.

C. Okan Sakar, S. O. Polat, Mete Katircioglu, and Yomi Kastro. Real-time prediction of online shoppers' purchasing intention using multilayer perceptron and lstm recurrent neural networks. *Neural Computing and Applications*, 31(10):6893–6908, 2019.